
Data Definition

- SQL DDL allows database objects such as schemas, domains, tables, views, and indexes to be created and destroyed.

- Main SQL DDL statements are:

CREATE SCHEMA DROP SCHEMA

CREATE/ALTER DOMAIN DROP DOMAIN

CREATE/ALTER TABLE DROP TABLE

CREATE VIEW DROP VIEW

- Many DBMSs also provide:

CREATE INDEX DROP INDEX

- Relations and other database objects exist in an environment.

- Each environment contains one or more catalogs, and each catalog consists of set of schemas.

- Schema is named collection of related database objects.

- Objects in a schema can be tables, views, domains, assertions, collations, translations, and character sets. All have same owner.

CREATE SCHEMA

CREATE SCHEMA [Name |
AUTHORIZATION CreatorId]

DROP SCHEMA Name [RESTRICT | CASCADE]

- With **RESTRICT** (default), schema must be empty or operation fails.

- With **CASCADE**, operation cascades to drop all objects associated with schema in order defined above. If any of these operations fail, DROP SCHEMA fails.

CREATE TABLE

CREATE TABLE TableName

{(colName dataType [NOT NULL] [UNIQUE]

[DEFAULT defaultOption]

[CHECK searchCondition] [...])

[PRIMARY KEY (listOfColumns),]

{[UNIQUE (listOfColumns),] [...]}

{[FOREIGN KEY (listOfFKColumns)

REFERENCES ParentTableName [(listOfCKColumns)],

[ON UPDATE referentialAction]

[ON DELETE referentialAction] [...])

{[CHECK (searchCondition)] [...]} }

- Creates a table with one or more columns of the specified dataType.

- With NOT NULL, system rejects any attempt to insert a null in the column.

- Can specify a DEFAULT value for the column.

- Primary keys should always be specified as NOT NULL.

- FOREIGN KEY clause specifies FK along with the referential action.

CREATE TABLE PropertyForRent (

propertyNo PNumber NOT NULL,

rooms PRooms NOT NULL DEFAULT 4,

rent PRent NOT NULL, DEFAULT 600,

ownerNo OwnerNumber NOT NULL,

staffNo StaffNumber

Constraint StaffNotHandlingTooMuch

branchNo BranchNumber NOT NULL,

PRIMARY KEY (propertyNo),

FOREIGN KEY (staffNo) REFERENCES Staff

ON DELETE SET NULL ON UPDATE CASCADE);

ALTER TABLE

- Add a new column to a table.
- Drop a column from a table.
- Add a new table constraint.
- Drop a table constraint.
- Set a default for a column.
- Drop a default for a column.

Change Staff table by removing default of 'Assistant' for position column and setting default for sex column to female ('F').

```
ALTER TABLE Staff
ALTER position DROP DEFAULT;
ALTER TABLE Staff
ALTER sex SET DEFAULT 'F';
```

Remove constraint from PropertyForRent that staff are not allowed to handle more than 100 properties at a time. Add new column to Client table.

```
ALTER TABLE PropertyForRent
DROP CONSTRAINT StaffNotHandlingTooMuch;
ALTER TABLE Client
ADD prefNoRooms PRooms;
```

DROP TABLE

DROP TABLE TableName [RESTRICT | CASCADE]

e.g. DROP TABLE PropertyForRent;

- Removes named table and all rows within it.
- With RESTRICT, if any other objects depend for their existence on continued existence of this table, SQL does not allow request.
- With CASCADE, SQL drops all dependent objects (and objects dependent on these objects).

Views

Dynamic result of one or more relational operations operating on base relations to produce another relation.

•Virtual relation that does not necessarily actually exist in the database but is produced upon request, at time of request.

- Contents of a view are defined as a query on one or more base relations.
- With view resolution, any operations on view are automatically translated into operations on relations from which it is derived.
- With view materialization, the view is stored as a temporary table, which is maintained as the underlying base tables are updated.

SQL - CREATE VIEW

```
CREATE VIEW ViewName [ (newColumnName [...]) ]
```

```
AS subselect
```

```
[WITH [CASCADED | LOCAL] CHECK OPTION]
```

- Can assign a name to each column in view.
- If list of column names is specified, it must have same number of items as number of columns produced by subselect.
- If omitted, each column takes name of corresponding column in subselect.

- List must be specified if there is any ambiguity in a column name.
- The subselect is known as the defining query.
- WITH CHECK OPTION ensures that if a row fails to satisfy WHERE clause of defining query, it is not added to underlying base table.
- Need SELECT privilege on all tables referenced in subselect and USAGE privilege on any domains used in referenced columns.

Create Horizontal View

Create view so that manager at branch B003 can only see details for staff who work in his or her office.

```
CREATE VIEW Manager3Staff
AS SELECT *
FROM Staff
WHERE branchNo = 'B003';
```

Create view of staff details at branch B003 excluding salaries.

```
CREATE VIEW Staff3
AS SELECT staffNo, fName, lName, position, sex
FROM Staff
WHERE branchNo = 'B003';
```

Grouped and Joined Views

Create view of staff who manage properties for rent, including branch number they work at, staff number, and number of properties they manage.

```
CREATE VIEW StaffPropCnt (branchNo, staffNo, cnt)
AS SELECT s.branchNo, s.staffNo, COUNT(*)
FROM Staff s, PropertyForRent p
WHERE s.staffNo = p.staffNo
GROUP BY s.branchNo, s.staffNo;
```

SQL - DROP VIEW

```
DROP VIEW ViewName [RESTRICT | CASCADE]
```

- Causes definition of view to be deleted from database.
- For example:
DROP VIEW Manager3Staff;

- With CASCADE, all related dependent objects are deleted; i.e. any views defined on view being dropped.

- With RESTRICT (default), if any other objects depend for their existence on continued existence of view being dropped, command is rejected.

View Resolution

Count number of properties managed by each member at branch B003.

```
SELECT staffNo, cnt
FROM StaffPropCnt
WHERE branchNo = 'B003'
ORDER BY staffNo;
```

(a) View column names in SELECT list are translated into their corresponding column names in the defining query:

```
SELECT s.staffNo As staffNo, COUNT(*) As cnt
```

(b) View names in FROM are replaced with corresponding FROM lists of defining query:

```
FROM Staff s, PropertyForRent p
```

(c) WHERE from user query is combined with WHERE of defining query using AND:

WHERE s.staffNo = p.staffNo AND branchNo = 'B003'

(d) GROUP BY and HAVING clauses copied from defining query:

GROUP BY s.branchNo, s.staffNo

(e) ORDER BY copied from query with view column name translated into defining query column name

ORDER BY s.staffNo

(f) Final merged query is now executed to produce the result:

SELECT s.staffNo AS staffNo, COUNT(*) AS cnt

FROM Staff s, PropertyForRent p

WHERE s.staffNo = p.staffNo AND
branchNo = 'B003'

GROUP BY s.branchNo, s.staffNo

ORDER BY s.staffNo;

Restrictions on Views

SQL imposes several restrictions on creation and use of views.

(a) If column in view is based on an aggregate function:

- Column may appear only in SELECT and ORDER BY clauses of queries that access view.
- Column may not be used in WHERE nor be an argument to an aggregate function in any query based on view.

•For example, following query would fail:

SELECT COUNT(cnt)
FROM StaffPropCnt;

•Similarly, following query would also fail:

SELECT *
FROM StaffPropCnt
WHERE cnt > 2;

(b) Grouped view may never be joined with a base table or a view.

•For example, StaffPropCnt view is a grouped view, so any attempt to join this view with another table or view fails.

View Updatability

•All updates to base table reflected in all views that encompass base table.

•Similarly, may expect that if view is updated then base table(s) will reflect change.

•However, consider again view StaffPropCnt.

•If we tried to insert record showing that at branch B003, SG5 manages 2 properties:

INSERT INTO StaffPropCnt
VALUES ('B003', 'SG5', 2);

•Have to insert 2 records into PropertyForRent showing which properties SG5 manages. However, do not know which properties they are; i.e. do not know primary keys!

•If change definition of view and replace count with actual property numbers:

CREATE VIEW StaffPropList (branchNo,
staffNo, propertyNo)
AS SELECT s.branchNo, s.staffNo, p.propertyNo
FROM Staff s, PropertyForRent p
WHERE s.staffNo = p.staffNo;

•Now try to insert the record:

INSERT INTO StaffPropList
VALUES ('B003', 'SG5', 'PG19');

- Still problem, because in PropertyForRent all columns except postcode/staffNo are not allowed nulls.

- However, have no way of giving remaining non-null columns values.

- ISO specifies that a view is updatable if and only if:

- DISTINCT is not specified.
- Every element in SELECT list of defining query is a column name and no column appears more than once.

- FROM clause specifies only one table, excluding any views based on a join, union, intersection or difference.

- No nested SELECT referencing outer table.

- No GROUP BY or HAVING clause.

- Also, every row added through view must not violate integrity constraints of base table.

For view to be updatable, DBMS must be able to trace any row or column back to its row or column in the source table.

WITH CHECK OPTION

- Rows exist in a view because they satisfy WHERE condition of defining query.

- If a row changes and no longer satisfies condition, it disappears from the view.

- New rows appear within view when insert/update on view cause them to satisfy WHERE condition.

- Rows that enter or leave a view are called migrating rows.

- WITH CHECK OPTION prohibits a row migrating out of the view.

- LOCAL/CASCADED apply to view hierarchies.

- With LOCAL, any row insert/update on view and any view directly or indirectly defined on this view must not cause row to disappear from view unless row also disappears from derived view/table.

- With CASCADED (default), any row insert/ update on this view and on any view directly or indirectly defined on this view must not cause row to disappear from the view.

```
CREATE VIEW Manager3Staff
AS SELECT *
FROM Staff
WHERE branchNo = 'B003'
WITH CHECK OPTION;
```

- Cannot update branch number of row B003 to B002 as this would cause row to migrate from view.

- Also cannot insert a row into view with a branch number that does not equal B003.

- Now consider the following:

```
CREATE VIEW LowSalary
AS SELECT * FROM Staff WHERE salary > 9000;
CREATE VIEW HighSalary
AS SELECT * FROM LowSalary
WHERE salary > 10000
WITH LOCAL CHECK OPTION;
CREATE VIEW Manager3Staff
AS SELECT * FROM HighSalary
WHERE branchNo = 'B003';
```

```
UPDATE Manager3Staff
SET salary = 9500
```

WHERE staffNo = 'SG37';

- This update would fail: although update would cause row to disappear from HighSalary, row would not disappear from LowSalary.
- However, if update tried to set salary to 8000, update would succeed as row would no longer be part of LowSalary.

- If HighSalary had specified WITH CASCADED CHECK OPTION, setting salary to 9500 or 8000 would be rejected because row would disappear from HighSalary.
- To prevent anomalies like this, each view should be created using WITH CASCADED CHECK OPTION.

Advantages of Views

- Data independence
- Currency
- Improved security
- Reduced complexity
- Convenience
- Customization
- Data integrity

Disadvantages of Views

- Update restriction
- Structure restriction
- Performance

View Materialization

- View resolution mechanism may be slow, particularly if view is accessed frequently.
- View materialization stores view as temporary table when view is first queried.
- Thereafter, queries based on materialized view can be faster than recomputing view each time.
- Difficulty is maintaining the currency of view while base table(s) are being updated.
- View maintenance aims to apply only those changes necessary to keep view current.
- Consider following view:

```
CREATE VIEW StaffPropRent(staffNo)
AS SELECT DISTINCT staffNo
   FROM PropertyForRent
   WHERE branchNo = 'B003' AND
      rent > 400;
```

- If insert row into PropertyForRent with rent £400 then view would be unchanged.
- If insert row for property PG24 at branch B003 with staffNo = SG19 and rent = 550, then row would appear in materialized view.
- If insert row for property PG54 at branch B003 with staffNo = SG37 and rent = 450, then no new row would need to be added to materialized view.
- If delete property PG24, row should be deleted from materialized view.
- If delete property PG54, then row for PG37 should not be deleted (because of existing property PG21).

Transactions

- SQL defines transaction model based on COMMIT and ROLLBACK.
- Transaction is logical unit of work with one or more SQL statements guaranteed to be atomic with respect to recovery.
- An SQL transaction automatically begins with a transaction-initiating SQL statement (e.g., SELECT, INSERT).

- Changes made by transaction are not visible to other concurrently executing transactions until transaction completes.

- Transaction can complete in one of four ways:

- COMMIT ends transaction successfully, making changes permanent.
- ROLLBACK aborts transaction, backing out any changes made by transaction.
- For programmatic SQL, successful program termination ends final transaction successfully, even if COMMIT has not been executed.
- For programmatic SQL, abnormal program end aborts transaction.

- New transaction starts with next transaction-initiating statement.

- SQL transactions cannot be nested.

- SET TRANSACTION configures transaction:

```
SET TRANSACTION  
[READ ONLY | READ WRITE] |  
[ISOLATION LEVEL READ UNCOMMITTED |  
READ COMMITTED|REPEATABLE READ |SERIALIZABLE ]
```

Immediate and Deferred Integrity Constraints

- Do not always want constraints to be checked immediately, but instead at transaction commit.

- Constraint may be defined as INITIALLY IMMEDIATE or INITIALLY DEFERRED, indicating mode the constraint assumes at start of each transaction.

- In former case, also possible to specify whether mode can be changed subsequently using qualifier [NOT] DEFERRABLE.

- Default mode is INITIALLY IMMEDIATE.

- SET CONSTRAINTS statement used to set mode for specified constraints for current transaction:

```
SET CONSTRAINTS  
{ALL | constraintName [, . . . ]}  
{DEFERRED ; IMMEDIATE}
```